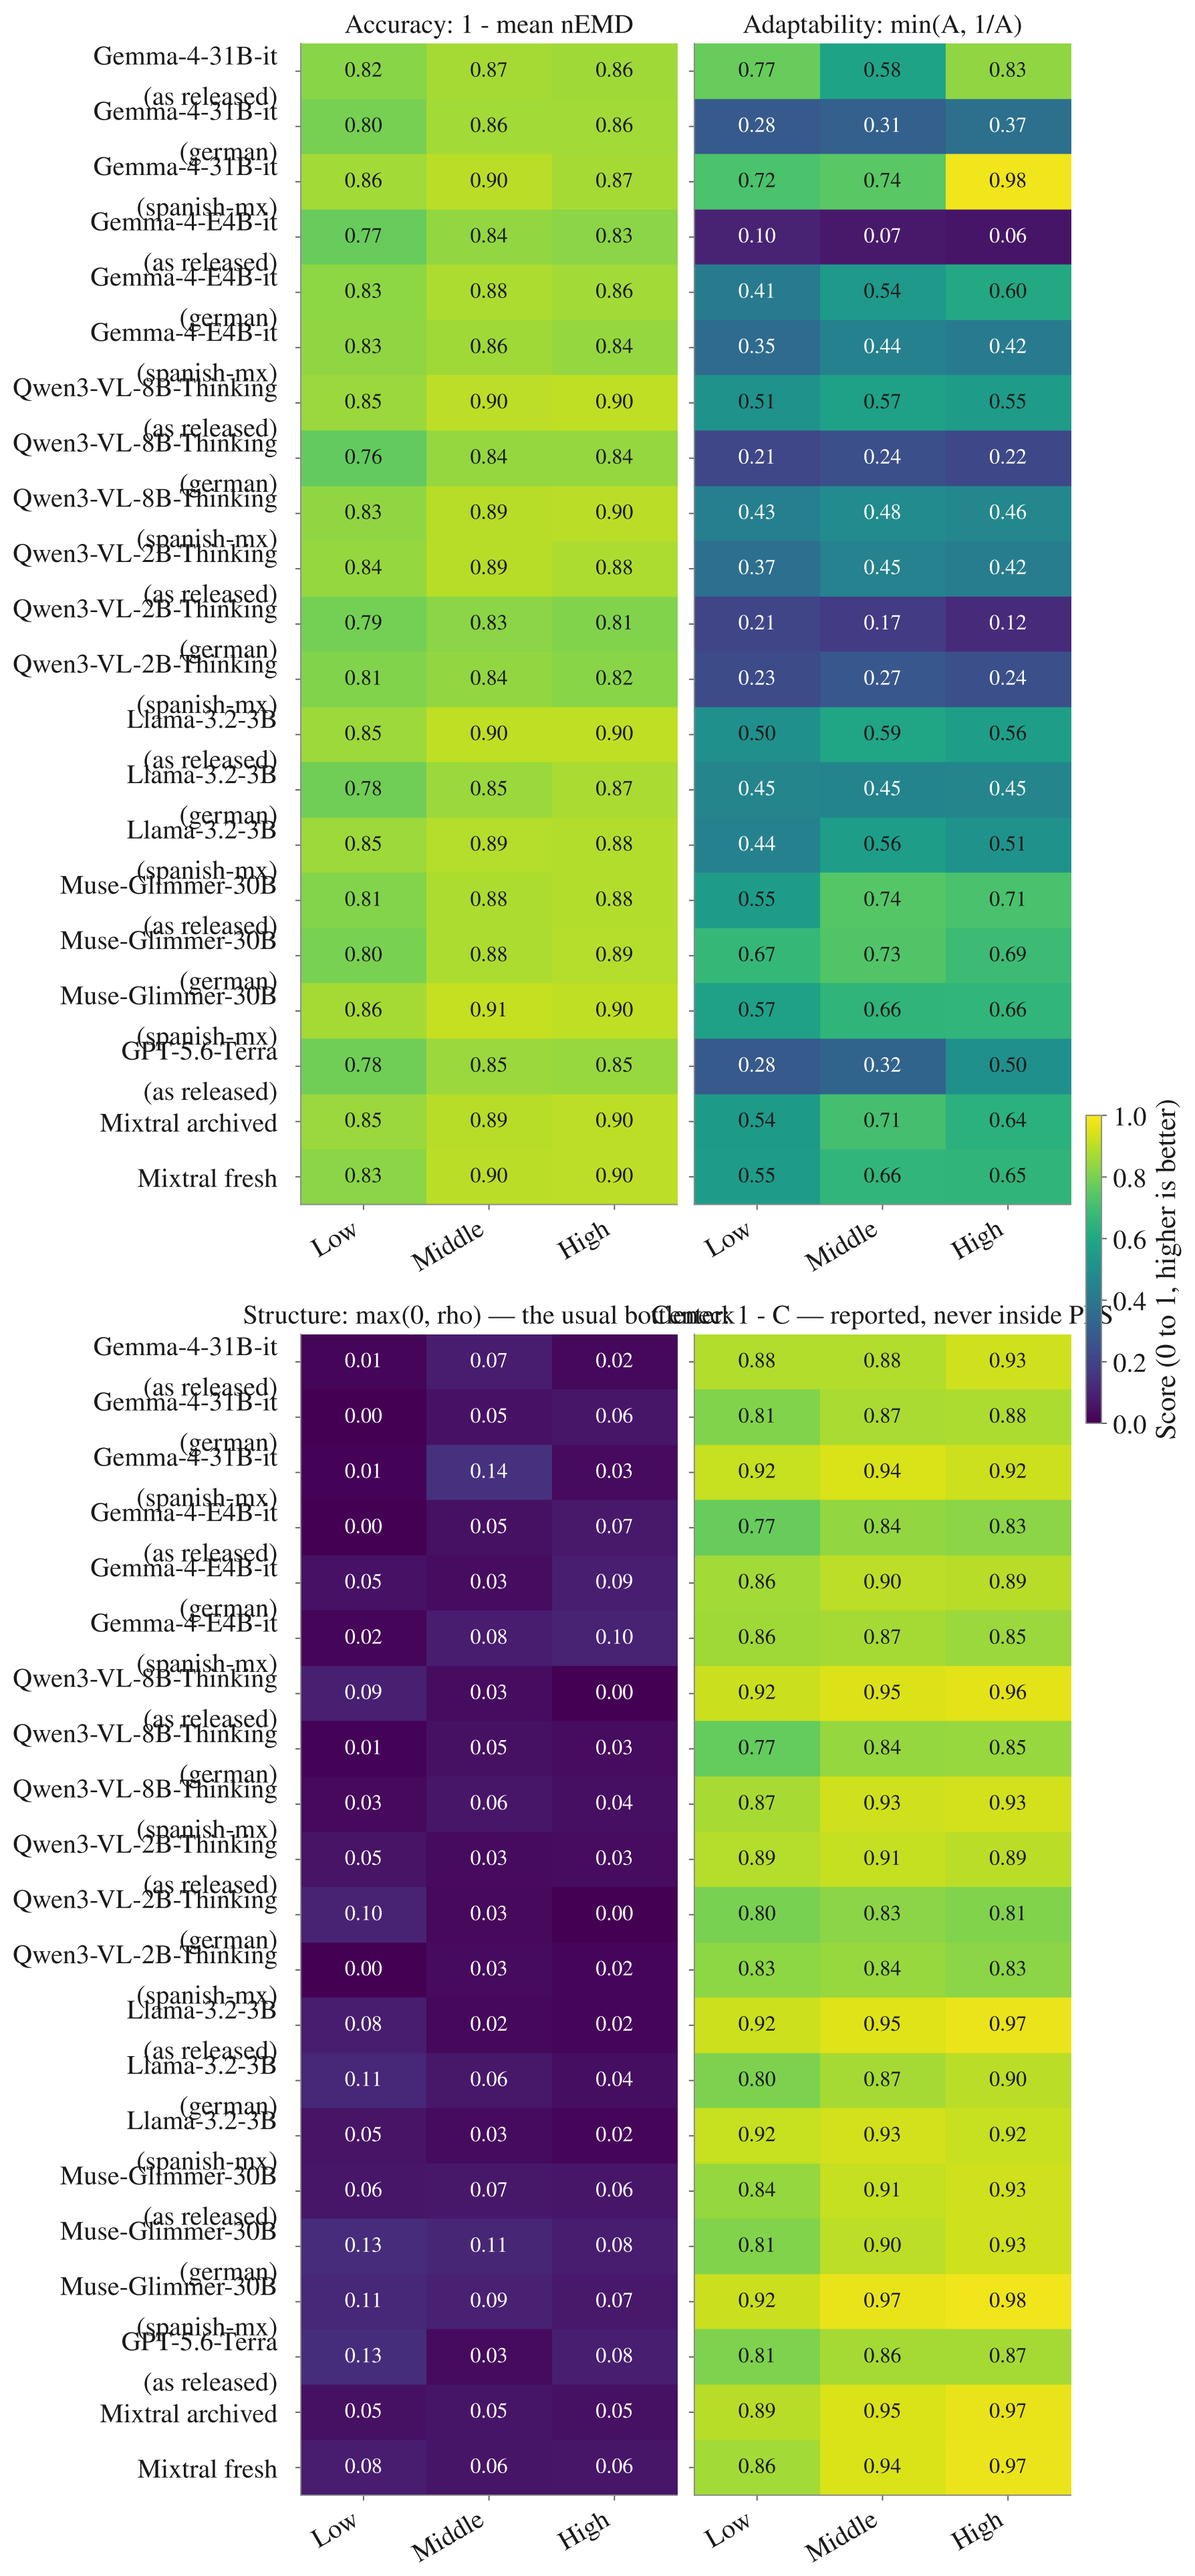

Accuracy: 1 - mean nEMD
Adaptability: min(A, 1/A)
Structure: max(0, rho) — the usual bottleneck
Clarity: 1 - C — reported, never inside P
Gemma-4-31B-it (as released)
Gemma-4-31B-it (german)
Gemma-4-31B-it (spanish-mx)
Gemma-4-E4B-it (as released)
Gemma-4-E4B-it (german)
Gemma-4-E4B-it (spanish-mx)
Qwen3-VL-8B-Thinking (as released)
Qwen3-VL-8B-Thinking (german)
Qwen3-VL-8B-Thinking (spanish-mx)
Qwen3-VL-2B-Thinking (as released)
Qwen3-VL-2B-Thinking (german)
Qwen3-VL-2B-Thinking (spanish-mx)
Llama-3.2-3B (as released)
Llama-3.2-3B (german)
Llama-3.2-3B (spanish-mx)
Muse-Glimmer-30B (as released)
Muse-Glimmer-30B (german)
Muse-Glimmer-30B (spanish-mx)
GPT-5.6-Terra (as released)
Mixtral archived
Mixtral fresh
Low
Middle
High
Score (0 to 1, higher is better)
1.0
0.8
0.6
0.4
0.2
0.0