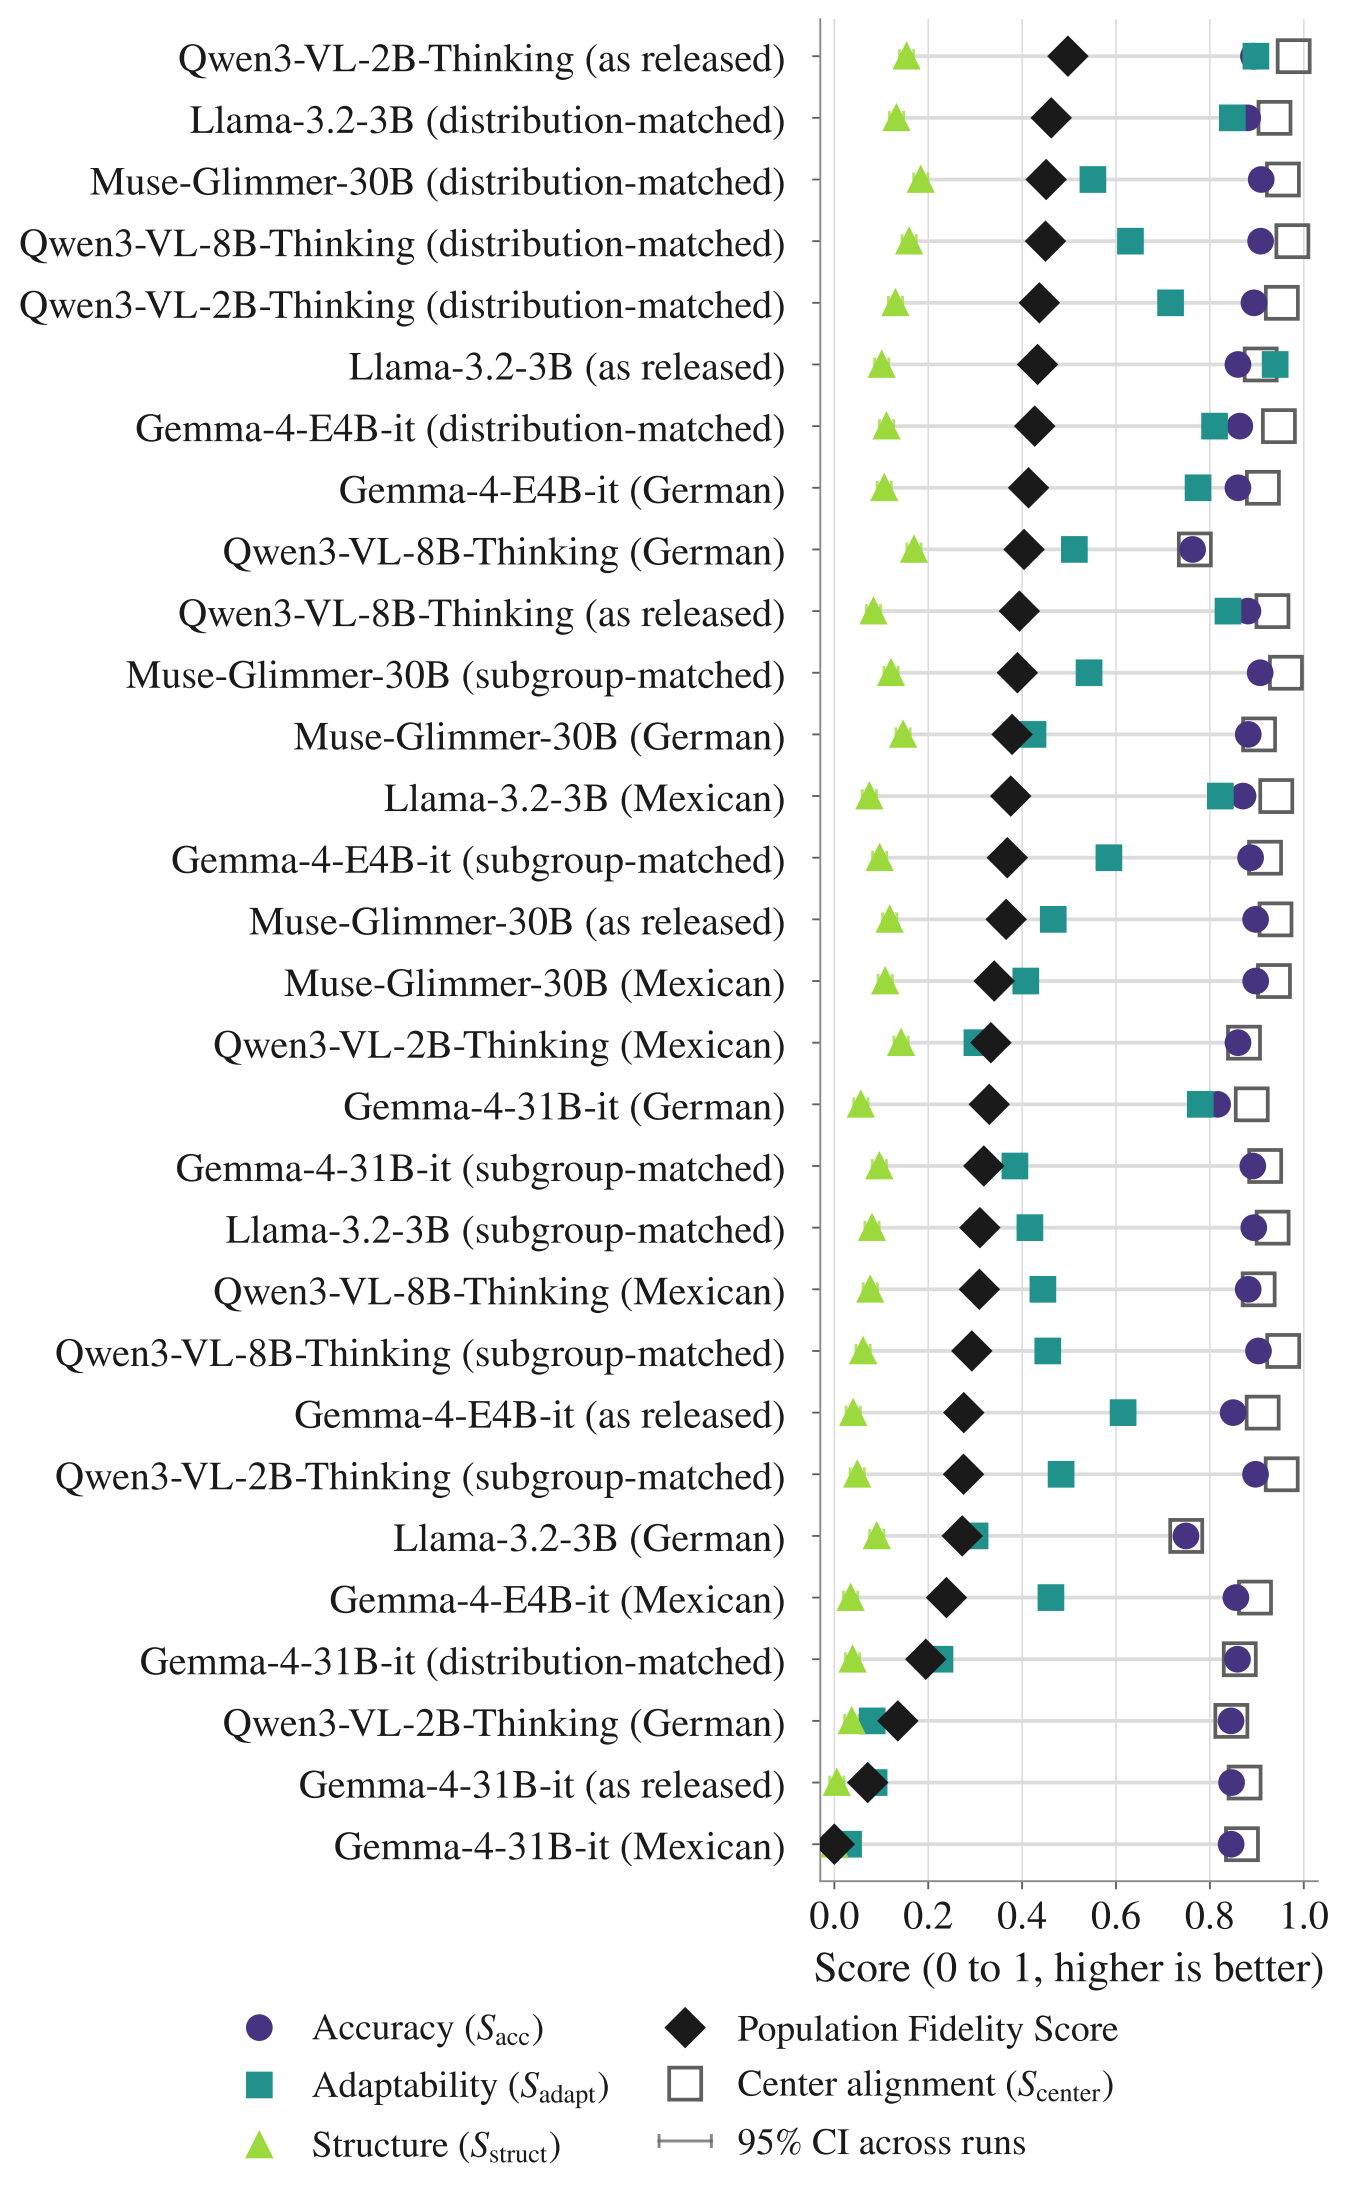
Qwen3-VL-2B-Thinking (as released)
Llama-3.2-3B (distribution-matched)
Muse-Glimmer-30B (distribution-matched)
Qwen3-VL-8B-Thinking (distribution-matched)
Qwen3-VL-2B-Thinking (distribution-matched)
Llama-3.2-3B (as released)
Gemma-4-E4B-it (distribution-matched)
Gemma-4-E4B-it (German)
Qwen3-VL-8B-Thinking (German)
Qwen3-VL-8B-Thinking (as released)
Muse-Glimmer-30B (subgroup-matched)
Muse-Glimmer-30B (German)
Llama-3.2-3B (Mexican)
Gemma-4-E4B-it (subgroup-matched)
Muse-Glimmer-30B (as released)
Muse-Glimmer-30B (Mexican)
Qwen3-VL-2B-Thinking (Mexican)
Gemma-4-31B-it (German)
Gemma-4-31B-it (subgroup-matched)
Llama-3.2-3B (subgroup-matched)
Qwen3-VL-8B-Thinking (Mexican)
Qwen3-VL-8B-Thinking (subgroup-matched)
Gemma-4-E4B-it (as released)
Qwen3-VL-2B-Thinking (subgroup-matched)
Llama-3.2-3B (German)
Gemma-4-E4B-it (Mexican)
Gemma-4-31B-it (distribution-matched)
Qwen3-VL-2B-Thinking (German)
Gemma-4-31B-it (as released)
Gemma-4-31B-it (Mexican)
0.0
0.2
0.4
0.6
0.8
1.0
Score (0 to 1, higher is better)
Accuracy (S_{acc})
Adaptability (S_{adapt})
Structure (S_{struct})
Population Fidelity Score
Center alignment (S_{center})
95% CI across runs